

Separate EBS (data) volume for storing/decoupling the IoT Data

Implementation

Dennis Catharina Johannes Kuijs

June 3, 2025

Contents

1. Context	2
2. Research based on feedback	3
2.1. Mount EBS volume on the Docker volumes directory	3
3. Implementation in the CI/CD pipeline	7
3.1. Creating and Mounting the EBS data volume	7
3.1.1. GitHub Actions workflow	7
3.1.2. Provision Host script	8
3.1.3. CloudFormation template	14
3.2. Adding CloudWatch metrics and alarms for the EBS data volume	15
3.2.1. CloudFormation Template	15
3.3. Adding support for automatic snapshot creation of the EBS data volume	20
3.3.1. Provision Host Script	20
3.3.2. CloudFormation template	22
3.4. Adding support for automatic attaching and detaching the EBS data volume	23
3.4.1. Provision Host script	23
3.4.2. CloudFormation Template	26
4. Improved implementation based on feedback	30
4.1. Add counter when attaching volume	30
4.2. Stop Docker before unmounting/detaching the volume	31
4.3. Wait until EC2 instance is created before attaching/mounting the volume	31
4.4. SSM Documents for each account instead of host	34
4.5. Improve check if filesystem exists	35
4.6. Problems with auto-reloader.conf	36
4.7. Attaching volume within the provision host script	37
4.8. Move provisioning DLM Policy to create-ec2 stack	38
4.9. Move provisioning EBS data volume to create-ec2 stack	39
5. Additional changes after second feedback round	41
5.1. Changing variable names and comments	41
5.2. Improving the filesystem check in the SSM document	43
5.3. Replace contents of existing volume	44
6. Tests	46

1. Context

This document provides a detailed overview how I implemented the creation and mounting of the separate `EBS` data volume within the `CI/CD` workflow. It also outlines the decisions I made and the challenges encountered throughout the development process.

2. Research based on feedback

After completing my initial research, I shared my findings with several team members. Following their feedback, I continued my research on this topic before proceeding with the implementation of my prototype in the `CI/CD` workflow.

2.1. Mount EBS volume on the Docker volumes directory

One of the team members suggested mounting the default `Docker` volumes directory (`/var/lib/docker/volumes`) to the newly created `EBS` data volume, rather than creating and mounting a custom directory.

I began investigating this option by mounting the separate `EBS` data volume to this directory using the following command:

```
sudo mount /dev/sdf /var/lib/docker/volumes
```

This works without any issues, when I checked the volume mount points using the following command.

```
sudo lsblk -f
```

It shows that the volume is mounted on the default `Docker` location (`/var/lib/docker/volumes`)

```
nvme1n1    xfs      a0bd2b05-85cf-4cb3-b79d-608d2db96ae6    99.2G    1% /var/lib/docker/volumes
```

Figure 1: EBS data volume is mounted to the `/var/lib/docker/volumes` directory

After starting OpenRemote with the default `Docker Compose` file, everything booted up properly without any issues.

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS
735d913f65f9	openremote/proxy:latest	"/entrypoint.sh run"	About a minute ago	Up 54 seconds (healthy)
5a9242762406	openremote/manager:latest	"/bin/sh -c 'java \$0...'"	About a minute ago	Up About a minute (healthy)
585f4a2ac2a0	openremote/keycloak:latest	"/bin/sh -c '/opt/ke...'"	About a minute ago	Up About a minute (healthy)
c0e5153edadd	openremote/postgresql:latest	"/or-entrypoint.sh p..."	About a minute ago	Up About a minute (healthy)

Figure 2: OpenRemote Docker containers are healthy

However, after attaching the `EBS` data volume to another `EC2` instance running OpenRemote, I encountered permission errors with the `PostgreSQL` container once again.

```

-----
pg_ctl: another server might be running; trying to start server anyway
waiting for server to start...2025-04-08 11:46:16.416 CEST [19] FATAL:  data directory "/var/lib/postgresql/data" has wrong ownership
2025-04-08 11:46:16.416 CEST [19] HINT:  The server must be started by the user that owns the data directory.
stopped waiting
pg_ctl: could not start server
Examine the log output.

```

Figure 3: PostgreSQL container has wrong ownership

Based on the insights of my initial research, I knew that this issue could be resolved by setting the `PGDATA` environment variable in the `Docker Compose` file.

Since the `EBS` volume is an external block device, this step is necessary for `Docker` to properly access the data. It's not possible to `chown` the directory to both the `postgres` and `root` users simultaneously, which makes specifying the `PGDATA` variable essential.

After adding the `PGDATA` variable the `Docker Compose` file looks like this.

```

# OpenRemote v3
#
# Profile that runs the stack by default on https://localhost using a self-signed SSL
↪ certificate,
# but optionally on https://$OR_HOSTNAME with an auto generated SSL certificate from
↪ Letsencrypt.
#
# It is configured to use the AWS logging driver.
#
volumes:
  proxy-data:
  manager-data:
  postgresql-data:

services:
  proxy:
    image: openremote/proxy:${PROXY_VERSION:-latest}
    restart: always
    depends_on:
      manager:
        condition: service_healthy
    ports:
      - 80:80 # Needed for SSL generation using letsencrypt
      - ${OR_SSL_PORT:-443}:443
      - 8883:8883
      - 127.0.0.1:8404:8404 # Localhost metrics access
    volumes:
      - proxy-data:/deployment
    environment:
      LE_EMAIL: ${OR_EMAIL_ADMIN:-}
      DOMAINNAME: ${OR_HOSTNAME:-localhost}
      DOMAINNAMES: ${OR_ADDITIONAL_HOSTNAMES:-}
      # USE A CUSTOM PROXY CONFIG - COPY FROM
      ↪ https://raw.githubusercontent.com/openremote/proxy/main/haproxy.cfg

```

```
# HAPROXY_CONFIG: '/data/proxy/haproxy.cfg'

postgresql:
  restart: always
  image: openremote/postgresql:${POSTGRES_VERSION:-latest}
  shm_size: 128mb
  volumes:
    - postgresql-data:/var/lib/postgresql/data
    - manager-data:/storage
  environment:
    PGDATA: /var/lib/postgresql/data/postgres

keycloak:
  restart: always
  image: openremote/keycloak:${KEYCLOAK_VERSION:-latest}
  depends_on:
    postgresql:
      condition: service_healthy
  volumes:
    - ./deployment:/deployment
  environment:
    KEYCLOAK_ADMIN_PASSWORD: ${OR_ADMIN_PASSWORD:-secret}
    KC_HOSTNAME: ${OR_HOSTNAME:-localhost}
    KC_HOSTNAME_PORT: ${OR_SSL_PORT:--1}

manager:
  # privileged: true
  restart: always
  image: openremote/manager:${MANAGER_VERSION:-latest}
  depends_on:
    keycloak:
      condition: service_healthy
  ports:
    - 127.0.0.1:8405:8405 # Localhost metrics access
  environment:
    OR_SETUP_TYPE:
    OR_ADMIN_PASSWORD:
    OR_SETUP_RUN_ON_RESTART:
    OR_EMAIL_HOST:
    OR_EMAIL_USER:
    OR_EMAIL_PASSWORD:
    OR_EMAIL_X_HEADERS:
    OR_EMAIL_FROM:
    OR_EMAIL_ADMIN:
    OR_METRICS_ENABLED: ${OR_METRICS_ENABLED:-true}
    OR_HOSTNAME: ${OR_HOSTNAME:-localhost}
    OR_ADDITIONAL_HOSTNAMES:
    OR_SSL_PORT: ${OR_SSL_PORT:--1}
    OR_DEV_MODE: ${OR_DEV_MODE:-false}

# The following variables will configure the demo
OR_FORECAST_SOLAR_API_KEY:
OR_OPEN_WEATHER_API_APP_ID:
```

```
OR_SETUP_IMPORT_DEMO_AGENT_KNX:  
OR_SETUP_IMPORT_DEMO_AGENT_VELBUS:  
volumes:  
- manager-data:/storage
```

With this setup, the `EBS` data volume can now easily be attached to other `EC2` instances as long as the `PGDATA` variable is configured on both the original and target machine. Additionally, the `Docker Compose` file becomes much simpler, only the `PGDATA` variable needs to be configured, eliminating the need to define different volume paths for each individual container.

3. Implementation in the CI/CD pipeline

In this section, I explain how I implemented my prototype into the existing `CI/CD` workflow on `Github Actions`.

It will be divided into the following topics:

- Creating and Mounting the `EBS` data volume
- Adding CloudWatch metrics and alarms for the `EBS` data volume
- Adding support for automatic snapshot creation of the `EBS` data volume
- Adding support for automatic attaching and detaching the `EBS` data volume

3.1. Creating and Mounting the EBS data volume

3.1.1. GitHub Actions workflow

I start my implementation in the `Github Actions` workflow file. In this file the steps for executing the `CI/CD` workflow are defined. The workflow is triggered on `workflow_dispatch`, meaning it runs on-demand without the need for a pull request or code push.

```
on:
  workflow_dispatch:
```

I added two additional input variables to this file: `DATA_DISK_SIZE` and `SNAPSHOT_ID`.

The `DATA_DISK_SIZE` variable allows you to specify the desired size of `EBS` data volume. By default, it is set to `16`, matching the size of the `root` device.

```
DATA_DISK_SIZE:
  description: 'Override EC2 data EBS volume size (GB)'
  type: string
  default: '16'
  required: false
```

The `SNAPSHOT_ID` variable allows you to specify a `snapshot`, enabling you to create a volume based off existing data. When this variable is specified the `DATA_DISK_SIZE` parameter is ignored. Instead, the volume will be provisioned with the same amount of storage that was assigned before `snapshot` creation.


```

SNAPSHOT_ID:
  description: 'Create EBS data volume based on snapshot'
  type: string
  required: false

```

Next, I added the input variables to the `.env` section in the `provision host` step. This ensures that the variable values can be accessed by referencing the input section at the top, as show below:

```

env:
  ACCOUNT_NAME: ${ github.event.inputs.ACCOUNT_NAME }
  HOST: ${ github.event.inputs.HOST }
  INSTANCE_TYPE: ${ github.event.inputs.INSTANCE_TYPE }
  ROOT_DISK_SIZE: ${ github.event.inputs.ROOT_DISK_SIZE }
  DATA_DISK_SIZE: ${ github.event.inputs.DATA_DISK_SIZE }
  SNAPSHOT_ID: ${ github.event.inputs.SNAPSHOT_ID }
  ELASTIC_IP: ${ github.event.inputs.ELASTIC_IP }
  PROVISION_S3_BUCKET: ${ github.event.inputs.PROVISION_S3_BUCKET }
  AWS_ACCESS_KEY_ID: ${ secrets._TEMP_AWS_ACCESS_KEY_ID }
  AWS_SECRET_ACCESS_KEY: ${ secrets._TEMP_AWS_SECRET_ACCESS_KEY }
  AWS_ROLE_NAME: ${ secrets._TEMP_AWS_ROLE_NAME }
  ENABLE_METRICS: ${ github.event.inputs.ENABLE_METRICS }

```

Finally, I passed the newly created variables to the `provision host` script. This ensures that the script can access the variable values and execute its logic based on them.

```

.ci_cd/aws/provision_host.sh "$ACCOUNT_NAME" "$HOST" "$INSTANCE_TYPE" "$ROOT_DISK_SIZE"
↩ "$DATA_DISK_SIZE" "$SNAPSHOT_ID" "$ELASTIC_IP" "$PROVISION_S3_BUCKET" "$ENABLE_METRICS"

```

3.1.2. Provision Host script

In the `provision host` script, I modified the order of the variables passed from the `workflow` to the script. In bash, you can reference each variable based on the order in which they are passed.

```

AWS_ACCOUNT_NAME=${1,,}
HOST=${2,,}
INSTANCE_TYPE=${3,,}
ROOT_DISK_SIZE=${4,,}
DATA_DISK_SIZE=${5,,}
SNAPSHOT_ID=${6,,}
ELASTIC_IP=${7,,}
PROVISION_S3_BUCKET=${8,,}
ENABLE_METRICS=${9,,}
WAIT_FOR_STACK=${10,,}

```

Next, I created the `EBS_STACK_NAME` variable, which generates a unique name for the `CloudFormation` stack by combining the `STACK_NAME` with a predefined text string. The `STACK_NAME` itself is created

from the `HOST` variable, where all dots in the `hostname` are replaced with a hyphen. With this approach, the `CloudFormation` stack names are unique for every `host`.

It is crucial that the `EBS` volume is created in the same `availability zone` as the `EC2` instance, otherwise, the volume cannot be attached to the instance. To ensure this, I first investigated how the `EC2` instance is assigned to a specific `availability zone`.

First, the `SUBNET_NUMBER` variable is set to a random integer between `1` and `3`. There are `3` different `public subnets` and this approach randomly selects one of them. Each subnet is located in a different `availability zone` (`1a`, `1b` or `1c`). The subnet name is then generated using the `SUBNET_NUMBER` variable and a predefined text string.

```
SUBNET_NUMBER=$(( $RANDOM % 3 + 1 ))
SUBNETNAME="or-subnet-public-$SUBNET_NUMBER"
```

I still needed the exact `availability zone` name that the `EC2` instance will use. The script already includes a line that retrieves the `AvailabilityZoneId` based on the `SUBNET_NAME` variable.

However, this ID cannot be used to create the `EBS` data volume, as the volume requires the `availability zone` name. To resolve this, I added the following line to the script to retrieve the name.

```
SUBNET_AZ=$(aws ec2 describe-subnets --filters Name=tag:Name,Values=$SUBNETNAME --query
↪ "Subnets[0].AvailabilityZone" --output text $ACCOUNT_PROFILE 2>/dev/null)
```

After setting the `availability zone` name to the `SUBNET_AZ` variable the script can start creating the `EBS` data volume. The volume creation is handled by a separate `CloudFormation` template to ensure that the `EBS` data volume will not be affected by updates to the `EC2` instance template.

Before creating the `EBS` data volume, the script checks if the stack not already exists.

```
STATUS=$(aws cloudformation describe-stacks --stack-name $EBS_STACK_NAME --query
↪ "Stacks[0].StackStatus" --output text 2>/dev/null)

if [ -n "$STATUS" ] && [ "$STATUS" != 'DELETE_COMPLETE' ]; then
    echo "Stack already exists for this host '$HOST' current status is '$STATUS'"
    EBS_STACK_ID=$(aws cloudformation describe-stacks --stack-name $EBS_STACK_NAME --query
↪ Stacks[0].StackId --output text 2>/dev/null)
else
```

If the stack exists, the `EBS_STACK_ID` variable will be set with the `Stack ID` for future reference. Otherwise, the `EBS` data volume is created.

Before creating the volume, the script checks if the `CloudFormation` template exists in one of the specified directories. If the template is not found, the script will exit and throw an error.

```

if [ -f "${awsDir}cloudformation-create-ebs-volume.yml" ]; then
    EBS_TEMPLATE_PATH="${awsDir}cloudformation-create-ebs-volume.yml"
elif [ -f ".ci_cd/aws/cloudformation-create-ebs-volume.yml" ]; then
    EBS_TEMPLATE_PATH=".ci_cd/aws/cloudformation-create-ebs-volume.yml"
    if [ -f "openremote/.ci_cd/aws/cloudformation-create-ebs-volume.yml" ]; then
        EBS_TEMPLATE_PATH="openremote/.ci_cd/aws/cloudformation-create-ebs-volume.yml"
    else
        echo "Cannot determine location of cloudformation-create-ebs-volume.yml"
        exit 1
    fi
fi

```

Next, the script configures the `HOST` , `AvailabilityZone` and `DiskSize` parameters that are specified in the `CloudFormation` template to create the volume. These values are either provided by the `workflow` inputs or generated earlier in the script.

```

PARAMS="ParameterKey=Host,ParameterValue=$HOST"
PARAMS="$PARAMS ParameterKey=AvailabilityZone,ParameterValue=$SUBNET_AZ"
PARAMS="$PARAMS ParameterKey=DiskSize,ParameterValue=$DATA_DISK_SIZE"

```

When the `SNAPSHOT_ID` variable is provided, this parameter will be configured to ensure that the volume is created based off an existing snapshot.

```

if [ -n "$SNAPSHOT_ID" ]; then
    PARAMS="$PARAMS ParameterKey=SnapshotId,ParameterValue='$SNAPSHOT_ID'"
fi

```

After configuring the parameters, the `CloudFormation` stack will be created with the following command. In this command I specify the stack name that was generated at the beginning of the script and pass the configured parameters.

```

EBS_STACK_ID=$(aws cloudformation create-stack --capabilities CAPABILITY_NAMED_IAM
↪ --stack-name $EBS_STACK_NAME --template-body file://$EBS_TEMPLATE_PATH --parameters
↪ $PARAMS --output text)

```

When the stack is successfully created, it returns the `Stack ID` , which is then be stored in the `EBS_STACK_ID` variable. The code below checks whether the stack creation command succeeded. If not, the script will throw an exit code and stop execution.

```

if [ $? -ne 0 ]; then
    echo "Create stack failed"
    exit 1
else
    echo "Create stack in progress"
fi

```

After the stack is successfully created, we need to check whether the creation was succesful or failed with an error. The code below retrieves the status from the `CloudFormation` stack based off the `Stack ID` stored in the previous step. As long as the status returns `CREATE_IN_PROGRESS` the stack is still being created.

The script checks the stack status every `30` seconds and stops when the status either returns `CREATE_COMPLETE` (indicating successful stack creation) or when the status is neither `CREATE_IN_PROGRESS` nor `CREATE_COMPLETE` (indicating stack creation failure).

```
echo "Waiting for stack to be created"
STATUS=$(aws cloudformation describe-stacks --stack-name $EBS_STACK_NAME --query
↪ "Stacks[?StackId=='$EBS_STACK_ID'].StackStatus" --output text 2>/dev/null)

while [[ "$STATUS" == 'CREATE_IN_PROGRESS' ]]; do
    echo "Stack creation is still in progress .. Sleeping 30 seconds"
    sleep 30
    STATUS=$(aws cloudformation describe-stacks --stack-name $EBS_STACK_NAME --query
↪ "Stacks[?StackId=='$EBS_STACK_ID'].StackStatus" --output text 2>/dev/null)
done

if [ "$STATUS" != 'CREATE_COMPLETE' ]; then
    echo "Stack creation has failed status is '$STATUS'"
    exit 1
else
    echo "Stack creation is complete"
fi
```

When the `EBS` data volume is succesfully created, it can be attached to the `EC2` instance. To attach the volume to the instance, you must specify a `device Name` such as `/dev/sda` , `/dev/sdb` etc.

It is not possible to automatically assign a `device name` when attaching the volume. You must specify a specific `device name` upfront. To achieve this, I configured a variable named `EBS_DEVICE_NAME` and set it to `/dev/sdf` as the designated `device name` .

HVM	/dev/sd[a-z] /dev/xvd[a-c][a-z] /dev/xvdd[a-x]	Differs by AMI /dev/sda1 or /dev/xvda	/dev/sd[b-z] /dev/xvdb[b-z] *	/dev/sd[b-e] /dev/sd[b-h] (h1.16xlarge) /dev/sd[b-y] (d2.8xlarge) /dev/sd[b-i] (i2.8xlarge) **
-----	--	--	-------------------------------------	--

Figure 4: The different virtualization methods that Amazon offers

As shown in the image above, for `EC2` instances that are using `HVM` as the virtualization method, it is recommended to choose a `device name` between `/dev/sd[b]` and `/dev/sd[z]` .

When the volume is successfully created it can be attached to the instance. Attaching the volume to the `EC2` instance is a technical process that involves several logical steps.

First, before the `EBS` data volume can be mounted there must of course be a running `EC2` instance. To check this the script retrieves the `Instance ID` and state from the `CloudFormation` stack that creates the `EC2` instance.

```
INSTANCE_ID=$(aws ec2 describe-instances --filters "Name=tag:Name,Values='$HOST'" --query  
↳ "Reservations[].Instances[?Tags[?Value=='$STACK_ID']].InstanceId" --output text  
↳ $ACCOUNT_PROFILE 2>/dev/null)  
INSTANCE_STATE=$(aws ec2 describe-instances --filters "Name=tag:Name,Values='$HOST'" --query  
↳ "Reservations[].Instances[?Tags[?Value=='$STACK_ID']].State.Name" --output text  
↳ $ACCOUNT_PROFILE 2>/dev/null)
```

The script checks for an `EC2` instance associated to the `CloudFormation` template by querying the specific `Stack ID` that was returned when this stack was successfully created.

To ensure that we always target the correct instance, a filter is applied to retrieve the instance with the name provided in the `CloudFormation` template. The script also retrieves the instance's status, as the volume can only be attached to an instance that is in the `running` state.

```
echo "Check if instance is available"  
count=0  
while [[ -z "$INSTANCE_ID" ]] && [[ "$INSTANCE_STATE" != 'running' ]] && [ $count -lt 30 ]; do  
    echo "Instance creation is still in progress .. Sleeping 30 seconds"  
    sleep 30  
    INSTANCE_ID=$(aws ec2 describe-instances --filters "Name=tag:Name,Values='$HOST'" --query  
↳ "Reservations[].Instances[?Tags[?Value=='$STACK_ID']].InstanceId" --output text  
↳ $ACCOUNT_PROFILE 2>/dev/null)  
    INSTANCE_STATE=$(aws ec2 describe-instances --filters "Name=tag:Name,Values='$HOST'"  
↳ --query "Reservations[].Instances[?Tags[?Value=='$STACK_ID']].State.Name" --output  
↳ text $ACCOUNT_PROFILE 2>/dev/null)  
    count=$((count+1))  
done
```

If the `Instance ID` is not found or the instance status is not `running`, the script waits for `30` seconds before retrying. Each attempt increments a counter, and the script continues checking as long as the counter remains below `30`.

This counter acts as a safeguard, in case the instance fails to launch successfully. If the counter exceeds `30`, the script will stop prevent running indefinitely.

```
if [ -z "$INSTANCE_ID" ] && [ "$INSTANCE_STATE" != 'running' ]; then  
    echo "Failed to provision instance"  
    exit 1  
fi
```

If the `Instance ID` cannot be retrieved or the instance state is not `running` after `30` attempts, the script will exit with an error status code.

Once the `Instance ID` is found and the instance is in a `running` state, the script attempts to attach the `EBS` data volume to the `EC2` instance. It's crucial that this step happens immediately after the instance becomes available, as several `cfn-scripts` begin running right after instance creation. One of these scripts automatically creates a filesystem on the volume and mounts it to the `/var/lib/docker/volumes` directory.

If the volume isn't attached in time, this step will fail, which then fails instance creation and therefore automatic rollbacks the `CloudFormation` stack.

After retrieving the `Instance ID` the script is trying to retrieve the `Volume ID` that belongs to the volume that was created by the `EBS` `CloudFormation` template.

```
VOLUME_ID=$(aws ec2 describe-volumes --filters "Name=tag:Name,Values='$HOST/data'" --query  
↪ "Volumes[?Tags[?Value=='$EBS_STACK_ID']].VolumeId" --output text $ACCOUNT_PROFILE  
↪ 2>/dev/null)
```

Once the `Volume ID` is found, the script attempts to attach the volume using the following command:

```
VOLUME=$(aws ec2 attach-volume --device $DEVICE_NAME --instance-id $INSTANCE_ID --volume-id  
↪ $VOLUME_ID --output text $ACCOUNT_PROFILE 2>/dev/null)
```

The command includes the configured `device name` from the previous step, the `Instance ID` to which the volume will be attached, and the ID of the volume itself. Immediately after, the status of the volume is retrieved to check if it has been successfully attached. If the volume is not attached, the script waits for `30` seconds before retrying.

```
while [[ "$STATUS" == 'attaching' ]]; do  
    echo "Volume is still attaching.. Sleeping 30 seconds"  
    sleep 30  
    STATUS=$(aws ec2 describe-volumes --filters "Name=tag:Name,Values='$HOST/data'" --query  
    ↪ "Volumes[?Tags[?Value=='$EBS_STACK_ID']].Attachments[].State" --output text  
    ↪ $ACCOUNT_PROFILE 2>/dev/null)  
done
```

When the status of the volume is `ATTACHED` the volume is attached successfully to the instance. If the status is anything else, the attachment failed, and the script will exit with an error status code.

```

if [ "$STATUS" != 'attached' ]; then
    echo "Volume attaching failed with status $STATUS"
    exit 1
else
    echo "Volume attaching is complete"
fi

```

After the `EBS` data volume is successfully attached to the `EC2` instance, the script waits for all `cfn-scripts` to complete execution on the instance. These scripts handle several tasks, including creating a filesystem on the attached volume and mounting it to `Docker's` default directory at `/var/lib/docker/volumes`. Additionally, the script updates the `/etc/fstab` file to ensure that the volume is automatically mounted on reboot using the device's `UUID` instead of the `device name`. This approach also prevents issues when attaching the same volume to another `EC2` instance.

```

prepare_volume:
  commands:
    01_mount_volume:
      command: !Sub |
        if [ -n "${SnapshotId}" ]; then
          sudo mount "${EBSDeviceName}" /var/lib/docker/volumes
        else
          sudo mkfs -t xfs "${EBSDeviceName}"
          sudo mount "${EBSDeviceName}" /var/lib/docker/volumes
        fi
    02_configure_fstab:
      command: !Sub |
        UUID=$(sudo blkid -o value -s UUID "${EBSDeviceName}")
        if [ -n $UUID ]; then
          sudo cp /etc/fstab /etc/fstab.orig
          sudo echo "UUID=$UUID /var/lib/docker/volumes xfs defaults,nofail 0 2" >>
↪ /etc/fstab
        else
          echo "Failed to create /etc/fstab entry. UUID is not found"
          exit 1
        fi

```

3.1.3. CloudFormation template

The `CloudFormation` template for creating the `EBS` volume looks like this:

```

AWSTemplateFormatVersion: '2010-09-09'
Description: 'Creates an EBS Volume for storing the IoT data.'
Parameters:
  Host:
    Description: The hostname of the machine where this volume is being attached.
    Type: String
  AvailabilityZone:

```

```

    Description: The AZ where the EBS volume needs to be created.
    Type: String
DiskSize:
    Description: Amount of storage you want to provision for this EBS volume.
    Type: Number
    Default: 16
SnapshotId:
    Description: Snapshot ID to create the EBS volume based of an existing Snapshot.
    Type: String
    Default: ""

Conditions:
    IsSnapshotProvided: !Not [!Equals [!Ref SnapshotId, ""]]

Resources:
    ORDataVolume:
    Type: AWS::EC2::Volume
    Properties:
    AvailabilityZone: !Ref AvailabilityZone
    Size: !Ref DiskSize
    VolumeType: gp3
    SnapshotId: !If [IsSnapshotProvided, !Ref SnapshotId, !Ref 'AWS::NoValue']
    Tags:
    - Key: Name
      Value: !Sub ${Host}/data

```

The script will create an `EBS` data volume based on the parameters that are passed from the `provision host` script. To easily identify each volume, a tag will be added with the host's name. If the condition `IsSnapshotProvided` is true, the `Snapshot ID` will be configured and the volume will be created based off an existing snapshot. Otherwise, an empty volume will be created.

3.2. Adding CloudWatch metrics and alarms for the EBS data volume

3.2.1. CloudFormation Template

To monitor the performance and health of the `EBS` data volume, I added `CloudWatch` metrics and alarms for this device. As part of the `cfn-scripts` the `Cloudwatch Agent` will be configured to collect various metrics from the `EC2` machine such as `CPU utilization`, `memory usage`, `disk usage` and more. The script looks like this:

```

{
  "agent": {
    "metrics_collection_interval": 300
  },
  "metrics": {
    "append_dimensions": {
      "InstanceId": "${aws:InstanceId}"
    }
  }
}

```



```
    },
    "metrics_collected": {
      "mem": {
        "measurement": [
          "mem_used_percent"
        ],
        "metrics_collection_interval": 900
      },
      "disk": {
        "drop_device": true,
        "measurement": [
          "used_percent"
        ],
        "resources": [
          "/",
          "/var/lib/docker/volumes"
        ],
        "metrics_collection_interval": 900
      }
    }
  },
  "logs": {
    "metrics_collected": {
      "prometheus": {
        "log_group_name": "Prometheus",
        "prometheus_config_path": "/opt/aws/amazon-cloudwatch-agent/var/prometheus.yaml",
        "emf_processor": {
          "metric_declaration_dedup": true,
          "metric_namespace": "CWAgent-Prometheus",
          "metric_unit": {
            "artemis_message_count": "Count",
            "artemis_messages_added": "Count",
            "or_rules_seconds_max": "Seconds",
            "or_rules_seconds_sum": "Seconds",
            "or_rules_seconds_count": "Count",
            "or_attributes_seconds_max": "Seconds",
            "or_attributes_seconds_sum": "Seconds",
            "or_attributes_seconds_count": "Count",
            "or_attributes_total": "Count",
            "or_provisioning_seconds_max": "Seconds",
            "or_provisioning_seconds_sum": "Seconds",
            "or_provisioning_seconds_count": "Count",
            "executor_pool_size_threads": "Count",
            "executor_pool_core_threads": "Count",
            "executor_pool_max_threads": "Count",
            "executor_seconds_count": "Count",
            "executor_seconds_sum": "Seconds",
            "haproxy_server_current_sessions": "Count",
            "haproxy_server_bytes_in_total": "Bytes",
            "haproxy_server_bytes_out_total": "Bytes",
            "haproxy_server_status": "Count",
            "haproxy_server_http_responses_total": "Count",
            "haproxy_server_max_session_rate": "Count/Second",
            "haproxy_server_total_time_average_seconds": "Seconds"
          }
        }
      }
    }
  }
}
```

```

    },
    "metric_declaration": [
      {
        "source_labels": [
          "job"
        ],
        "label_matcher": "^manager$",
        "dimensions": [
          "InstanceName"
        ],
      },
      {
        "metric_selectors": [
          "^or_rules_seconds_count$",
          "^or_rules_seconds_sum$",
          "^or_rules_seconds_max$",
          "^or_attributes_seconds_count$",
          "^or_attributes_seconds_sum$",
          "^or_attributes_seconds_max$",
          "^or_provisioning_seconds_count$",
          "^or_provisioning_seconds_sum$",
          "^or_provisioning_seconds_max$"
        ],
      },
    ],
    {
      "source_labels": [
        "job",
        "source"
      ],
      "label_matcher":
↪ "manager;(RulesEngine|AgentService|DefaultMQTTHandler|AssetResource|WebsocketClient)",
      "dimensions": [
        "InstanceName",
        "source"
      ],
    },
    {
      "metric_selectors": [
        "^or_attributes_total$"
      ],
    },
    {
      "source_labels": [
        "job",
        "name"
      ],
      "label_matcher": "^manager;ContainerExecutor$",
      "dimensions": [
        "InstanceName",
        "name"
      ],
    },
    {
      "metric_selectors": [

```

```
        "^executor_pool_",
        "^executor_seconds_count$",
        "^executor_seconds_sum$"
    ]
},
{
    "source_labels": [
        "job"
    ],
    "label_matcher": "^manager$",
    "dimensions": [
        [
            "InstanceName",
            "queue"
        ]
    ],
    "metric_selectors": [
        "^artemis_message_count$",
        "^artemis_messages_added$"
    ]
},
{
    "source_labels": [
        "job"
    ],
    "label_matcher": "^proxy$",
    "dimensions": [
        [
            "InstanceName",
            "proxy",
            "server"
        ]
    ],
    "metric_selectors": [
        "^haproxy_server_total_time_average_seconds$",
        "^haproxy_server_max_session_rate$",
        "^haproxy_server_bytes",
        "^haproxy_server_current_sessions$"
    ]
},
{
    "source_labels": [
        "job"
    ],
    "label_matcher": "^proxy$",
    "dimensions": [
        [
            "InstanceName",
            "proxy",
            "server",
            "code"
        ]
    ],
    "metric_selectors": [
```

```
    "^haproxy_server_http_responses_total$"
  ],
  {
    "source_labels": [
      "job"
    ],
    "label_matcher": "^proxy$",
    "dimensions": [
      "InstanceName",
      "proxy",
      "server",
      "state"
    ]
  },
  "metric_selectors": [
    "^haproxy_server_status$"
  ]
}
]
}
}
}
```

To monitor disk usage, I added the mount point for the new `EBS` data volume to the `resources` block. This ensures that the `CloudWatch Agent` retrieves metrics for this specific mount point.

```
"drop_device": true,  
"measurement": [  
  "used_percent"  
],  
"resources": [  
  "/",  
  "/var/lib/docker/volumes"  
],
```

Next, I added a new `CloudWatch Alarm` to the same `CloudFormation` template. This alarm is set to trigger if disk usage exceeds `90%` within a `one-hour` period. When the alarm is triggered, Amazon sends a notification to the configured `SNS` topic, which in turn sends an email alert to the topic's subscribers.

To connect the newly added metric to this alarm, I configured the `Dimensions` block with the required details, including the `Instance ID` , the `path` that is referring to the `mount point` in the metrics and the `filesystem type` .

```

DataDiskUtilizationAlarm:
  Type: AWS::CloudWatch::Alarm
  Condition: MetricsEnabled
  Properties:
    Namespace: CWAgent
    MetricName: disk_used_percent
    Statistic: Average
    Period: 3600
    EvaluationPeriods: 1
    ComparisonOperator: GreaterThanThreshold
    Threshold: 90
    AlarmActions:
      - !Ref SnsTopic
    OKActions:
      - !Ref SnsTopic
    Dimensions:
      - Name: InstanceId
        Value: !Ref EC2Instance
      - Name: path
        Value: /var/lib/docker/volumes
      - Name: fstype
        Value: xfs

```

3.3. Adding support for automatic snapshot creation of the EBS data volume

3.3.1. Provision Host Script

To ensure that the data on the `EBS` data volume is securely backed up, the `provision host` script creates an `Amazon Data Lifecycle Manager (DLM)` policy for automatic `snapshot` creation. This policy ensures that snapshots of the `EBS` data volume are automatically created at regular intervals. To maintain consistency throughout the script, I implemented this feature in the same way as the `EBS` data volume creation.

I began by setting the `DLM_STACK_NAME` variable to generate an unique `CloudFormation` stack name for this feature.

```
DLM_STACK_NAME="$STACK_NAME-dlm-ebs-snapshot-policy"
```

Next, the script checks if the `CloudFormation` stack for the `DLM` policy already exists, if it does, the `STACK_ID` variable is set for future reference and the script will continue with the next steps.

```

if [ -n "$STATUS" ] && [ "$STATUS" != 'DELETE_COMPLETE' ]; then
  echo "Stack already exists for this host '$HOST' current status is '$STATUS'"
  STACK_ID=$(aws cloudformation describe-stacks --stack-name $DLM_STACK_NAME --query "Stacks[0].StackId" -
  -output text 2>/dev/null)
else

```

If the `CloudFormation` stack does not exist, the script will attempt to create it by first searching for the `CloudFormation` template within the specified directories. If the `CloudFormation` template cannot be found, the script will exit with an error status code. Otherwise the stack creation process will proceed.

```
if [ -f "${awsDir}cloudformation-create-dlm-policy.yml" ]; then
    DLM_TEMPLATE_PATH="${awsDir}cloudformation-create-dlm-policy.yml"
elif [ -f ".ci_cd/aws/cloudformation-create-dlm-policy.yml" ]; then
    DLM_TEMPLATE_PATH=".ci_cd/aws/cloudformation-create-dlm-policy.yml"
elif [ -f "openremote/.ci_cd/aws/cloudformation-create-dlm-policy.yml" ]; then
    DLM_TEMPLATE_PATH="openremote/.ci_cd/aws/cloudformation-create-dlm-policy.yml"
else
    echo "Cannot determine location of cloudformation-create-dlm-policy.yml"
    exit 1
fi
```

Before the stack can be created, the script first checks if the required `IAM` role is already created in the AWS account. This role is required for the `DLM` policy to perform `snapshot` creation tasks on behalf of the `IAM` user.

If the role doesn't exist, the system will create the default role and set the `ROLE_ARN` variable with the `Amazon Resource Name` (ARN). This variable will then be passed to the `CloudFormation` template in the next step.

```
echo "Check if IAM Role exists"
ROLE_ARN=$(aws iam get-role --role-name AWSDataLifecycleManagerDefaultRole --query "Role.Arn"
↪ --output text $ACCOUNT_PROFILE)

if [ -z "$ROLE_ARN" ]; then
    ROLE=$(aws dlm create-default-role --resource-type snapshot)

    if [ $? -ne 0 ]; then
        echo "IAM Role creation has failed"
        exit 1
    else
        echo "IAM Role creation is complete"
    fi

    ROLE_ARN=$(aws iam get-role --role-name AWSDataLifecycleManagerDefaultRole --query
↪ "Role.Arn" --output text $ACCOUNT_PROFILE)
fi
```

After configuring the `ROLE_ARN` variable, the script sets the necessary variables for the `CloudFormation` template. The `DLM_DESCRIPTION` variable is constructed by combining the `HOST` variable and removing any periods, as they are not allowed in the description. Additionally, the `ROLE_ARN` retrieved in the previous step is set, along with the `EBS_STACK_ID`, which identifies the volume that the policy should target.

```
DLM_DESCRIPTION="OpenRemote-${HOST%.*}"
PARAMS="ParameterKey=PolicyDescription,ParameterValue='${DLM_DESCRIPTION}'"
PARAMS="$PARAMS ParameterKey=DLMExecutionRoleArn,ParameterValue='${ROLE_ARN}'"
PARAMS="$PARAMS ParameterKey=EBSStackId,ParameterValue='${EBS_STACK_ID}'"
```

Once the parameters are configured, the script attempts to create the `CloudFormation` stack. If the stack creation fails, the script will exit with an error code and stops execution.

```
STACK_ID=$(aws cloudformation create-stack --capabilities CAPABILITY_NAMED_IAM --stack-name
↪ $DLM_STACK_NAME --template-body file://$DLM_TEMPLATE_PATH --parameters $PARAMS --output
↪ text)

if [ $? -ne 0 ]; then
    echo "Create stack failed"
    exit 1
fi
```

After the `CloudFormation` stack is successfully created, the script checks its status every `30` seconds. If the status returns `CREATE_COMPLETE`, the stack was created successfully, and the script proceeds to the next steps. However, if the status is neither `CREATE_IN_PROGRESS` nor `CREATE_COMPLETE`, it indicates an error, and the script will exit with an error code and stop execution.

```
if [ "$WAIT_FOR_STACK" != 'false' ]; then
    # Wait for CloudFormation stack status to be CREATE_*
    echo "Waiting for stack to be created"
    STATUS=$(aws cloudformation describe-stacks --stack-name $DLM_STACK_NAME --query
↪ "Stacks[?StackId=='$STACK_ID'].StackStatus" --output text 2>/dev/null)

    while [ "$STATUS" == 'CREATE_IN_PROGRESS' ]; do
        echo "Stack creation is still in progress .. Sleeping 30 seconds"
        sleep 30
        STATUS=$(aws cloudformation describe-stacks --stack-name $DLM_STACK_NAME --query
↪ "Stacks[?StackId=='$STACK_ID'].StackStatus" --output text 2>/dev/null)
    done

    if [ "$STATUS" != 'CREATE_COMPLETE' ]; then
        echo "Stack creation has failed status is '$STATUS'"
        exit 1
    else
        echo "Stack creation is complete"
    fi
fi
```

3.3.2. CloudFormation template

The `CloudFormation` template for creating the `DLM` policy looks like this:

```

Parameters:
  PolicyDescription:
    Description: Lifecycle Policy Description
    Type: String
  DLMEExecutionRoleArn:
    Description: Role ARN for executing the DLM operations.
    Type: String
  EBSStackId:
    Description: EBS StackId of the volume that need to be targeted by this policy.
    Type: String

Resources:
  EBSPolicy:
    Type: AWS::DLM::LifecyclePolicy
    Properties:
      Description: !Ref PolicyDescription
      ExecutionRoleArn: !Ref DLMEExecutionRoleArn
      State: ENABLED
      PolicyDetails:
        PolicyLanguage: STANDARD
        PolicyType: EBS_SNAPSHOT_MANAGEMENT
        TargetTags:
          - Key: aws:cloudformation:stack-id
            Value: !Ref EBSStackId
        ResourceTypes:
          - VOLUME
        Schedules:
          - Name: Daily Backup
            CreateRule:
              Interval: 24
              IntervalUnit: HOURS
              Times:
                - '05:00'
            RetainRule:
              Count: 5

```

The script creates a `lifecycle policy` with a single schedule that creates a new `snapshot` every `24 hours` at `5 AM`. The `RetainRule` is configured to keep the `5` most recent snapshots, automatically deleting older ones beyond that limit. The policy is enabled immediately upon creation.

To ensure that snapshots are only created for the correct volume, the `TargetTags` parameter uses the `EBS_STACK_ID` value to identify the appropriate volume.

3.4. Adding support for automatic attaching and detaching the EBS data volume

3.4.1. Provision Host script

The final step in the process is to create two scripts for automatically `attaching` and `detaching` the `EBS` data volume. These scripts are stored in `Amazon Systems Manager` and can be executed either

manually through the AWS Management Console or via the AWS CLI .

To provision the SSM scripts, I reused the same logic applied when provisioning the EBS data volume and DLM policy in the provision host script.

First, I added a new variable at the beginning of the script to generate the CloudFormation stack name.

```
SSM_STACK_NAME="$STACK_NAME-ssm-attach-detach-documents"
```

Next, the script checks if the CloudFormation stack for generating the SSM documents already exists.

```
echo "Provisioning SSM Documents for attaching/detaching EBS Data volume"
STATUS=$(aws cloudformation describe-stacks --stack-name $SSM_STACK_NAME --query
↪ "Stacks[0].StackStatus" --output text 2>/dev/null)

if [ -n "$STATUS" ] && [ "$STATUS" != 'DELETE_COMPLETE' ]; then
    echo "Stack already exists for this host '$HOST' current status is '$STATUS'"
    STACK_ID=$(aws cloudformation describe-stacks --stack-name $SSM_STACK_NAME --query
↪ "Stacks[0].StackId" --output text 2>/dev/null)
else
```

If the stack already exists, the STACK_ID variable will be set with the Stack ID for future reference. Otherwise, the script continues with the stack creation process. The first step in this process is to verify whether the CloudFormation template for creating the SSM documents exists within the specified directories.

```
if [ -f "${awsDir}cloudformation-create-ssm-document" ]; then
    SSM_TEMPLATE_PATH="${awsDir}cloudformation-create-ssm-document.yml"
elif [ -f ".ci_cd/aws/cloudformation-create-ssm-document.yml" ]; then
    SSM_TEMPLATE_PATH=".ci_cd/aws/cloudformation-create-ssm-document.yml"
elif [ -f "openremote/.ci_cd/aws/cloudformation-create-ssm-document.yml" ]; then
    SSM_TEMPLATE_PATH="openremote/.ci_cd/aws/cloudformation-create-ssm-document.yml"
else
    echo "Cannot determine location of cloudformation-create-ssm-document.yml"
    exit 1
fi
```

If the template is not found, the script exits with an error code and stops execution. Otherwise, it proceeds to configure the required parameters for provisioning the SSM documents.

The CloudFormation template required several values such as the INSTANCE_ID , VOLUME_ID , DEVICE_NAME and HOST , which are used within the scripts to handle the attachment and detachment of the EBS volume.

```
PARAMS="ParameterKey=Host,ParameterValue='$HOST'"
PARAMS="$PARAMS ParameterKey=InstanceId,ParameterValue='$INSTANCE_ID'"
PARAMS="$PARAMS ParameterKey=VolumeId,ParameterValue='$VOLUME_ID'"
PARAMS="$PARAMS ParameterKey=EBSDeviceName,ParameterValue='$DEVICE_NAME'"
```

After setting the parameters, the script attempts to create the `CloudFormation` stack.

```
STACK_ID=$(aws cloudformation create-stack --capabilities CAPABILITY_NAMED_IAM --stack-name
↪ $SSM_STACK_NAME --template-body file://$SSM_TEMPLATE_PATH --parameters $PARAMS --output
↪ text)
```

If stack creation fails, the system will throw an error code and stops execution.

```
if [ $? -ne 0 ]; then
    echo "Create stack failed"
    exit 1
fi
```

After the stack is successfully created, the script waits for the stack creation process to complete if the `WAIT_FOR_STACK` variable is set to true. During this time, it checks the stack status every `30` seconds and returns a success once the status changes to `CREATE_COMPLETE`. If the status is neither `CREATE_IN_PROGRESS` nor `CREATE_COMPLETE`, it indicates that the stack creation has failed. In that case, the script exits with an error code and stops execution.

```
if [ "$WAIT_FOR_STACK" != 'false' ]; then
    # Wait for CloudFormation stack status to be CREATE_*
    echo "Waiting for stack to be created"
    STATUS=$(aws cloudformation describe-stacks --stack-name $SSM_STACK_NAME --query
↪ "Stacks[?StackId=='$STACK_ID'].StackStatus" --output text 2>/dev/null)

    while [[ "$STATUS" == 'CREATE_IN_PROGRESS' ]]; do
        echo "Stack creation is still in progress .. Sleeping 30 seconds"
        sleep 30
        STATUS=$(aws cloudformation describe-stacks --stack-name $SSM_STACK_NAME --query
↪ "Stacks[?StackId=='$STACK_ID'].StackStatus" --output text 2>/dev/null)
    done

    if [ "$STATUS" != 'CREATE_COMPLETE' ]; then
        echo "Stack creation has failed status is '$STATUS'"
        exit 1
    else
        echo "Stack creation is complete"
    fi
fi
```

3.4.2. CloudFormation Template

The `CloudFormation` template for provisioning the `SSM` documents looks like this:

```
AWSTemplateFormatVersion: '2010-09-09'
Description: 'Creates an SSM Document for attaching/detaching the EBS volume'

Parameters:
  Host:
    Description: FQDN for host.
    Type: String
  InstanceId:
    Description: InstanceId where the script needs to be executed.
    Type: String
  VolumeId:
    Description: VolumeId that needs to be attached/detached.
    Type: String
  EBSDeviceName:
    Description: EBS DeviceName where this volume is mounted on.
    Type: String

Resources:
  SSMDetachEBSDocument:
    Type: AWS::SSM::Document
    Properties:
      Content:
        schemaVersion: '2.2'
        description: 'Script for detaching the EBS volume'
        parameters:
          InstanceId:
            type: String
            description: InstanceId where the script needs to be executed.
            default: !Ref InstanceId
          VolumeId:
            type: String
            description: VolumeId that needs to be detached.
            default: !Ref VolumeId
          EBSDeviceName:
            type: String
            description: EBS DeviceName where this volume is mounted on.
            default: !Ref EBSDeviceName
        mainSteps:
          - name: RemoveFstabEntry
            action: aws:runShellScript
            inputs:
              runCommand:
                - |
                  UUID=$(sudo blkid -o value -s UUID {{ EBSDeviceName }})
                  if [ -n "$UUID" ]; then
                    cp /etc/fstab /etc/fstab.orig
                    sed -i '/UUID='$UUID'/d' /etc/fstab
                  else
                    echo "Failed to remove /etc/fstab entry. UUID is not found"
                    exit 1
```

```

        fi
    - name: StopDocker
      action: aws:runShellScript
      inputs:
        runCommand:
          - systemctl stop docker
    - name: UmountVolume
      action: aws:runShellScript
      inputs:
        runCommand:
          - umount {{ EBSDeviceName }}
    - name: DetachVolume
      action: aws:runShellScript
      inputs:
        runCommand:
          - |
            VOLUME=$(aws ec2 detach-volume --volume-id {{ VolumeId }})
            STATUS=$(aws ec2 describe-volumes --query "Volumes[?VolumeId=='{{
↪ VolumeId }}'].State" --output text 2>/dev/null)

            while [[ "$STATUS" == 'in-use' ]] do
              echo "Instance is still in-use .. Sleeping 30 seconds"
              sleep 30
              STATUS=$(aws ec2 describe-volumes --query "Volumes[?VolumeId=='{{
↪ VolumeId }}'].State" --output text 2>/dev/null)
            done

            if [ "$STATUS" != "available" ]; then
              echo "Failed to detach volume"
              exit 1
            fi
        fi

DocumentFormat: YAML
TargetType: /AWS::EC2::Instance
UpdateMethod: Replace
DocumentType: Command
Name: !Sub ${Host}_detach

SSMAttachEBSDocument:
  Type: AWS::SSM::Document
  Properties:
    Content:
      schemaVersion: '2.2'
      description: 'Script for attaching the EBS volume'
      parameters:
        InstanceId:
          type: String
          description: InstanceId where the script needs to be executed.
          default: !Ref InstanceId
        VolumeId:
          type: String
          description: VolumeId that needs to be attached.
          default: !Ref VolumeId
        EBSDeviceName:
          type: String

```

```

    description: EBS DeviceName where this volume needs to be mounted on.
    default: !Ref EBSDeviceName
mainSteps:
  - name: AttachVolume
    action: aws:runShellScript
    inputs:
      runCommand:
        - |
          VOLUME=$(aws ec2 attach-volume --device {{ EBSDeviceName }}
↪ --instance-id {{ InstanceId }} --volume-id {{ VolumeId }})
          STATUS=$(aws ec2 describe-volumes --query "Volumes[?VolumeId=='{{
↪ VolumeId }}'].Attachments[].State" --output text 2>/dev/null)

          while [[ "$STATUS" == 'attaching' ]] do
            echo "Volume is still attaching .. Sleeping 30 seconds"
            sleep 30
            STATUS=$(aws ec2 describe-volumes --query "Volumes[?VolumeId=='{{
↪ VolumeId }}'].Attachments[].State" --output text 2>/dev/null)
          done

          if [ "$STATUS" != 'attached' ]; then
            echo "Volume attaching failed with status $STATUS"
            exit 1
          else
            echo "Volume attaching is complete"
          fi

  - name: MountVolume
    action: aws:runShellScript
    inputs:
      runCommand:
        - mount {{ EBSDeviceName }} /var/lib/docker/volumes
  - name: AddFstabEntry
    action: aws:runShellScript
    inputs:
      runCommand:
        - |
          UUID=$(sudo blkid -o value -s UUID {{ EBSDeviceName }})
          if [ -n "$UUID" ]; then
            cp /etc/fstab /etc/fstab.orig
            echo "UUID=$UUID /var/lib/docker/volumes xfs defaults,nofail 0 2" >>
↪ /etc/fstab

          else
            echo "Failed to add /etc/fstab entry. UUID is not found"
            exit 1
          fi

  - name: StartDocker
    action: aws:runShellScript
    inputs:
      runCommand:
        - systemctl start docker
DocumentFormat: YAML
TargetType: /AWS::EC2::Instance
UpdateMethod: Replace

```

```
DocumentType: Command
Name: !Sub ${Host}_attach
```

It creates two separate documents, one for `attaching` and another for `detaching` the `EBS` data volume. When these documents are executed, `SSM` runs the commands defined in the `runCommand` block on the targeted `EC2` instance.

4. Improved implementation based on feedback

After my initial implementation, I reviewed it with a team member and received valuable feedback to improve it further.

4.1. Add counter when attaching volume

In the first version, I added logic in the `create-ssm` CloudFormation template to `attach` and `mount` the volume, and to create a filesystem if it doesn't already exist. However, since volume attachment can sometimes fail or take longer than expected, the while loop I used could potentially run indefinitely.

To address this issue, I added a counter to limit the number of attempts. With this change, the loop will exit after `30` attempts (approximately `900` seconds).

```
while [[ "$STATUS" == 'available' ]] && [ $count -lt 30 ]; do
  echo "Volume is still attaching .. Sleeping 30 seconds"
  sleep 30
  STATUS=$(aws ec2 describe-volumes --query "Volumes[?VolumeId=='${ VolumeId }'].State"
  ↪ --output text 2>/dev/null)
  count=$((count+1))
done
```

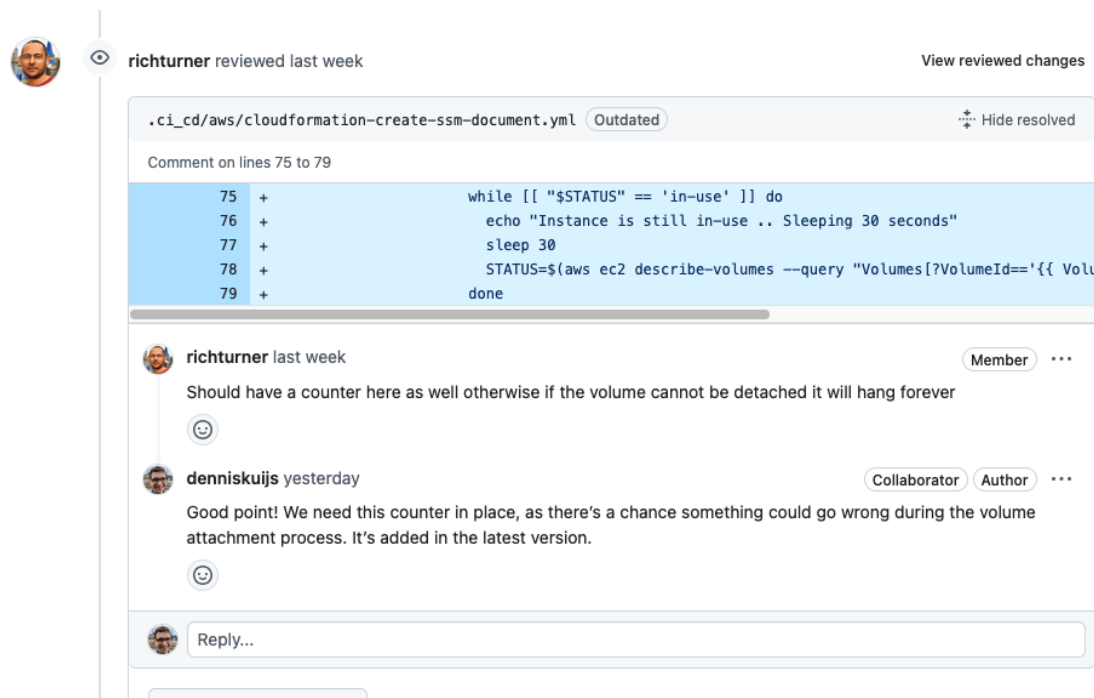


Figure 5: Rich proposed to add a counter when attaching the EBS data volume

4.2. Stop Docker before unmounting/detaching the volume

The second change I made was to add an extra command in the `create-ssm` `CloudFormation` template to stop `Docker` before `unmounting` and `detaching` the volume.

Previously, if the volume was `unmounted` and `detached` without stopping `Docker`, `OpenRemote` would lose access to the data, causing the containers become unhealthy. By stopping `Docker` first, `OpenRemote` is taken offline safely and can be restarted once the volume is reattached.

```
- name: StopDocker
  action: aws:runShellScript
  inputs:
    runCommand:
      - systemctl stop docker.socket docker.service
```

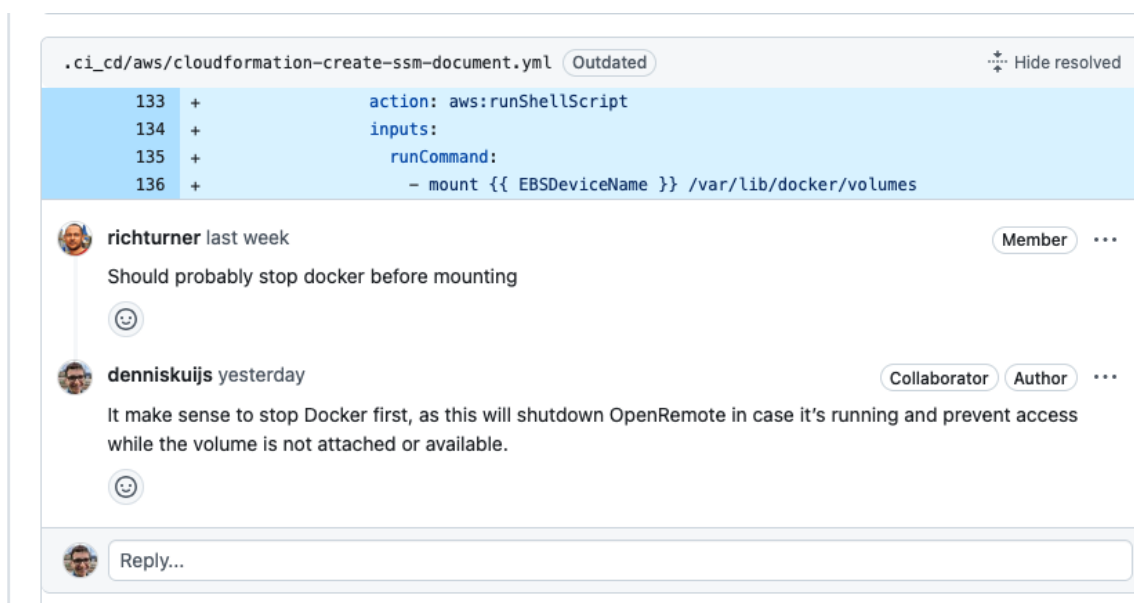


Figure 6: Rich proposed to stop the Docker service first before unmounting the EBS data volume

4.3. Wait until EC2 instance is created before attaching/mounting the volume

In my initial attempt, I added logic to provision the `EBS` data volume using a separate `CloudFormation` template. Once the instance was created by the `create-ec2` stack, I attached the volume immediately. The goal with this approach was to ensure the volume would be mounted in time before the `cfn-scripts` responsible for formatting and mounting the volume were executed.

However, the creation of the `EC2` instance is only considered complete once all `cfn-scripts` have been successfully executed and `cfn-signal` returns a `SUCCESS` status. Since volume creation and

mounting were part of the `cfn-scripts`, It wasn't possible to wait for the instance to be initialized beforehand. As a result, if the volume couldn't be attached or mounted successfully or in time, the `cfn-scripts` would fail, and the `CloudFormation` stack will be rolled back.

To address this issue, I moved the provisioning of the `EBS` data volume into the `create-ec2` stack. Since the volume must reside in the same `availability zone` as the instance, it is now created immediately after the instance is successfully launched.

```
EBSDataVolume:
  Type: AWS::EC2::Volume
  Properties:
    AvailabilityZone: !GetAtt EC2Instance.AvailabilityZone
    Size: !Ref DataDiskSize
    VolumeType: gp3
    SnapshotId: !If [SnapshotProvided, !Ref SnapshotId, !Ref 'AWS::NoValue']
  Tags:
    - Key: Name
      Value: !Sub ${Host}/data
```

After the instance is successfully created, I added logic in the `provision host` script to execute the `SSM` document responsible for `attaching` and `mounting` the `EBS` data volume. As part of this change, the previous logic for handling volume attachment and mounting was removed from the `cfn-scripts` section within the `create-ec2` `CloudFormation` template.

```
PARAMS="InstanceId=$INSTANCE_ID,VolumeId=$VOLUME_ID,DeviceName=$EBS_DEVICE_NAME"

COMMAND_ID=$(aws ssm send-command --document-name attach_volume --instance-ids $INSTANCE_ID
↪ --parameters $PARAMS --query "Command.CommandId" --output text $ACCOUNT_PROFILE
↪ 2>/dev/null)

if [ $? -ne 0 ]; then
  echo "Volume attaching/mounting failed"
  exit 1
fi

STATUS=$(aws ssm get-command-invocation --command-id $COMMAND_ID --instance-id $INSTANCE_ID
↪ --query "StatusDetails" --output text $ACCOUNT_PROFILE 2>/dev/null)

while [[ "$STATUS" == 'InProgress' ]]; do
  echo "Volume attaching/mounting is still in progress .. Sleeping 30 seconds"
  sleep 30
  STATUS=$(aws ssm get-command-invocation --command-id $COMMAND_ID --instance-id
↪ $INSTANCE_ID --query "StatusDetails" --output text $ACCOUNT_PROFILE 2>/dev/null)
done

if [ "$STATUS" != 'Success' ]; then
  echo "Volume attaching/mounting has failed status is '$STATUS'"
  exit 1
else
```

```
echo "Volume attaching/mounting is complete"
fi
```

With this approach, volume attachment and mounting only occur after the instance has been successfully created. This ensures the volume operations are only executed on a fully initialized instance, reducing the risks of errors during provisioning when for example the volume is not attached to the instance on time before executing the `cfn-scripts`.

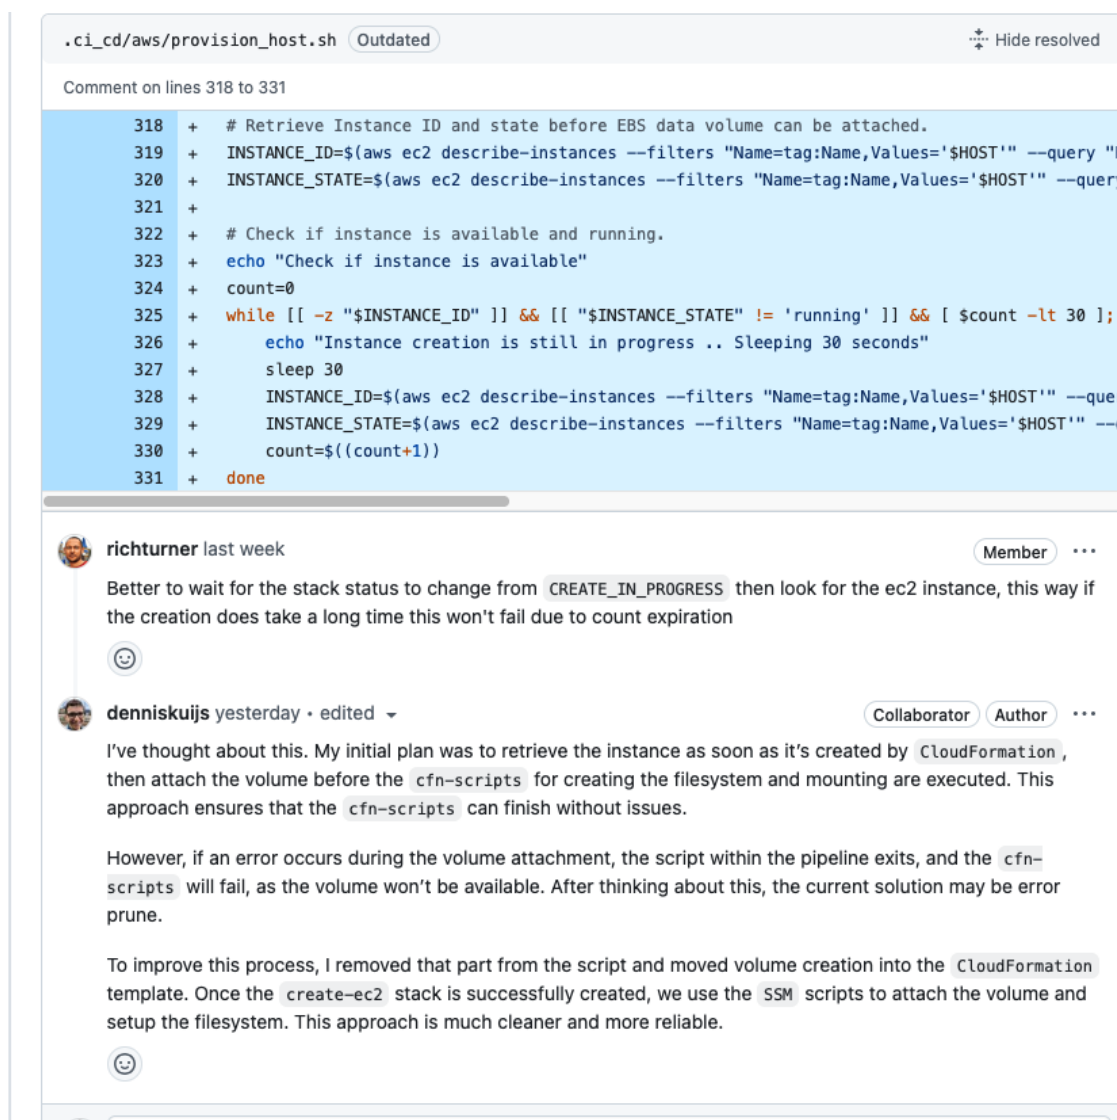


Figure 7: Rich proposed to check whether the EBS data volume is attached on time

4.4. SSM Documents for each account instead of host

When provisioning the `SSM` documents. In my first attempt, they will be created for each individual host. Since the documents are configurable with parameters and we need to specify on which instance they need to be executed. There is no reason to create this for every single host. Instead, provisioning them for each account is a better alternative. To achieve this, I moved this logic to the `provision account` script.

```
# Provision SSM Documents
if [ -f "${awsDir}cloudformation-create-ssm-documents.yml" ]; then
    TEMPLATE_PATH="${awsDir}cloudformation-create-ssm-documents.yml"
elif [ -f ".ci_cd/aws/cloudformation-create-ssm-documents.yml" ]; then
    TEMPLATE_PATH=".ci_cd/aws/cloudformation-create-ssm-documents.yml"
elif [ -f "openremote/.ci_cd/aws/cloudformation-create-ssm-documents.yml" ]; then
    TEMPLATE_PATH="openremote/.ci_cd/aws/cloudformation-create-ssm-documents.yml"
else
    echo "Cannot determine location of cloudformation-create-ssm-documents.yml"
    exit 1
fi

STACK_NAME=or-ebs-volume-ssm-documents

# Create SSM Documents for attaching/detaching EBS data volume in specified account
STACK_ID=$(aws cloudformation create-stack --capabilities CAPABILITY_NAMED_IAM --stack-name
↪ $STACK_NAME --template-body file://${TEMPLATE_PATH} --output text $ACCOUNT_PROFILE)

# Wait for CloudFormation stack status to be CREATE_*
echo "Waiting for stack to be created"
STATUS=$(aws cloudformation describe-stacks --stack-name $STACK_NAME --query
↪ "Stacks[0].StackStatus" --output text $ACCOUNT_PROFILE 2>/dev/null)

while [[ "$STATUS" == 'CREATE_IN_PROGRESS' ]]; do
    echo "Stack creation is still in progress .. Sleeping 30 seconds"
    sleep 30
    STATUS=$(aws cloudformation describe-stacks --stack-name $STACK_NAME --query
    ↪ "Stacks[0].StackStatus" --output text $ACCOUNT_PROFILE 2>/dev/null)
done

if [ "$STATUS" != 'CREATE_COMPLETE' ] && [ "$STATUS" != 'UPDATE_COMPLETE' ]; then
    echo "Stack creation has failed status is '$STATUS'" >&2
    exit 1
else
    echo "Stack creation is complete"
fi
```

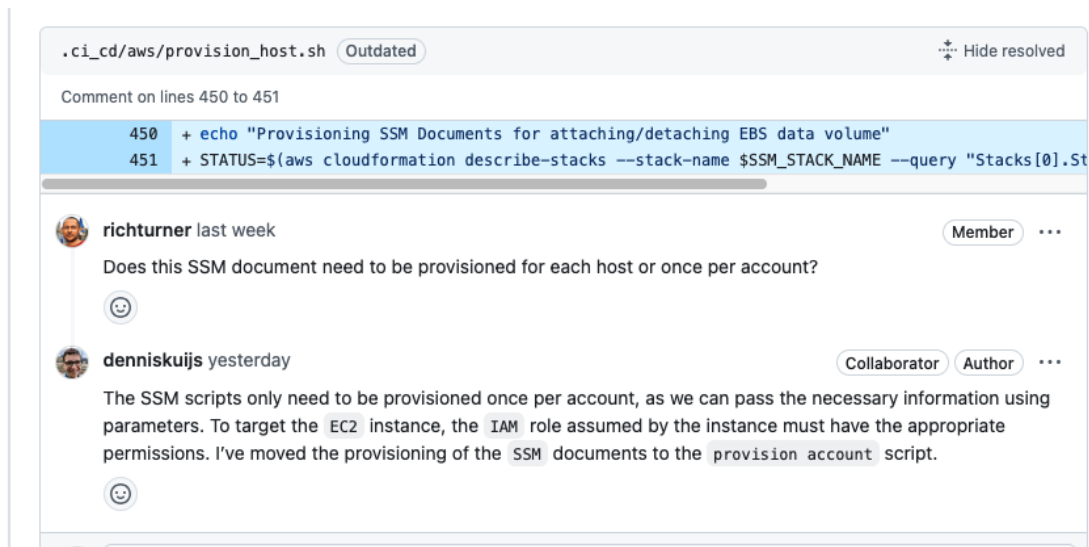


Figure 8: Rich proposed to create the SSM documents for each account instead of each host

4.5. Improve check if filesystem exists

In my initial implementation, I used the `snapshot` variable to determine whether the script should create a filesystem on the `EBS` data volume or simply mount it. When a `snapshot` is used, the volume already contains a filesystem, so it only needs to be mounted. Creating a new filesystem in this case would overwrite the existing one, resulting in data loss. Therefore, it's important to ensure that a filesystem is only created when no `snapshot` is provided.

After improving the implementation, I now use the `blkid` command to check for an existing filesystem. This command also retrieves the volume's `UUID` for generating the `/etc/fstab` entry.

Since the `ssm` commands are executed asynchronously, the system doesn't wait for previous commands to complete before continuing. However, retrieving the filesystem information is only possible once the volume is properly attached. To handle this, I added a while loop that waits for a filesystem to be detected before proceeding with the script.

```
count=0
while [[ -z "$FILESYSTEM" ]] && [ $count -lt 30 ]; do
    echo "Filesystem is currently not available .. Sleeping 30 seconds"
    sleep 30
    FILESYSTEM=$(blkid -o value -s TYPE {{ DeviceName }})
    count=$((count+1))
done

if [ -z "$FILESYSTEM" ]; then
    mkfs -t xfs {{ DeviceName }}
```

```

mount {{ DeviceName }} /var/lib/docker/volumes
else
mount {{ DeviceName }} /var/lib/docker/volumes
fi

```

Currently, the implementation is not optimal. When creating the volume for the first time, there is no filesystem present, which causes the script to loop up to 30 times (approximately 900 seconds) before proceeding with creating a filesystem. This delay is too long, so I'm planning to improve the implementation to handle this more efficiently.

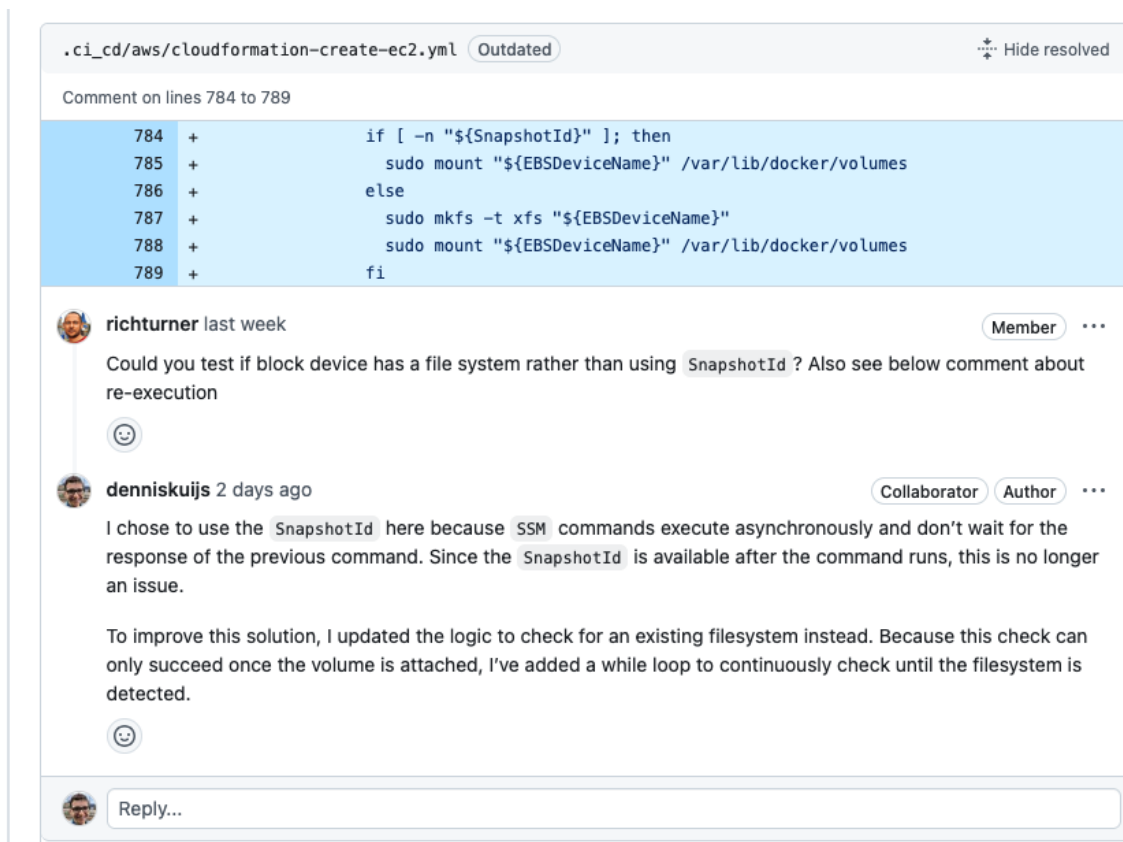


Figure 9: Rich proposed to change the filesystem check to look for an existing filesystem instead for an snapshot

4.6. Problems with auto-reloader.conf

Updating the CloudFormation stack re-executes the cfn-scripts , causing duplicate entries in /etc/fstab

```

services:
  systemd:
    cfn-hup:
      enabled: true
      ensureRunning: true
      files:
        - /etc/cfn/cfn-hup.conf
        - /etc/cfn/hooks.d/amazon-cloudwatch-agent-auto-reloader.conf

```

In my new implementation, I moved the volume attachment and mounting logic out of the `cfn-scripts` and into the `SSM` documents. This change ensures the logic isn't re-executed every time the instance is updated.

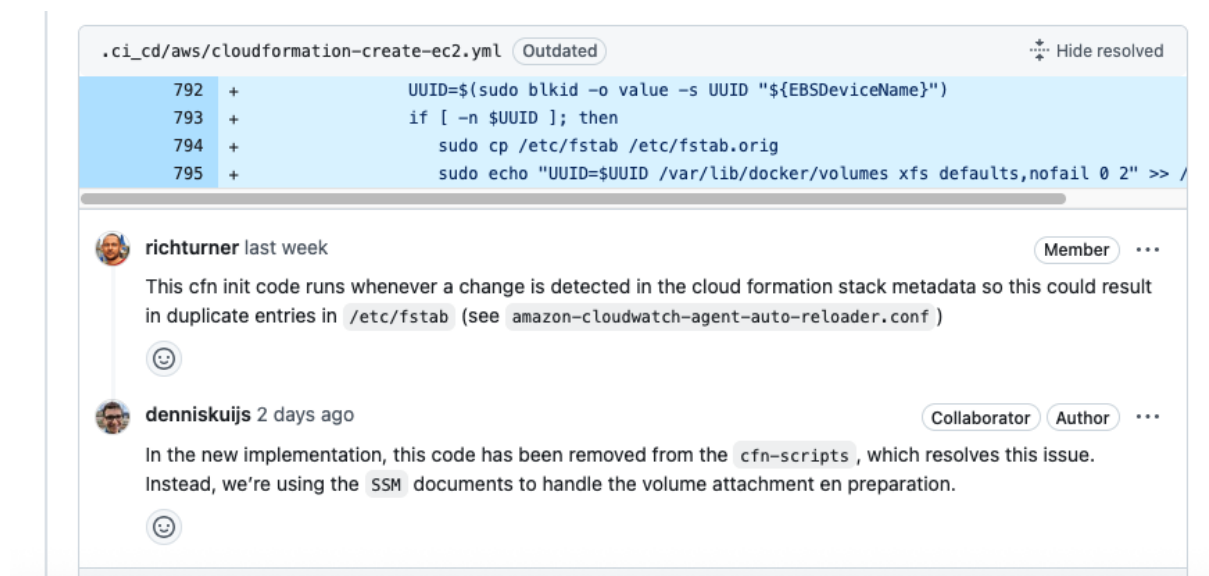


Figure 10: Rich proposed to take a look at the auto-reloader configuration as it might conflict with the current implementation

4.7. Attaching volume within the provision host script

Initially, I added the logic for attaching the volume in the `provision host` script. However, as previously mentioned, this approach could lead to errors if the volume wasn't created on time when the script ran. To resolve this, I moved the provisioning of the `EBS` data volume to the `create-ec2` CloudFormation template and removed it from the `provision host` script. Now, once the `instance` is successfully created and the `cfn-signal` returns a `SUCCESS` status, the volume is `attached` and `mounted` using `SSM` documents.

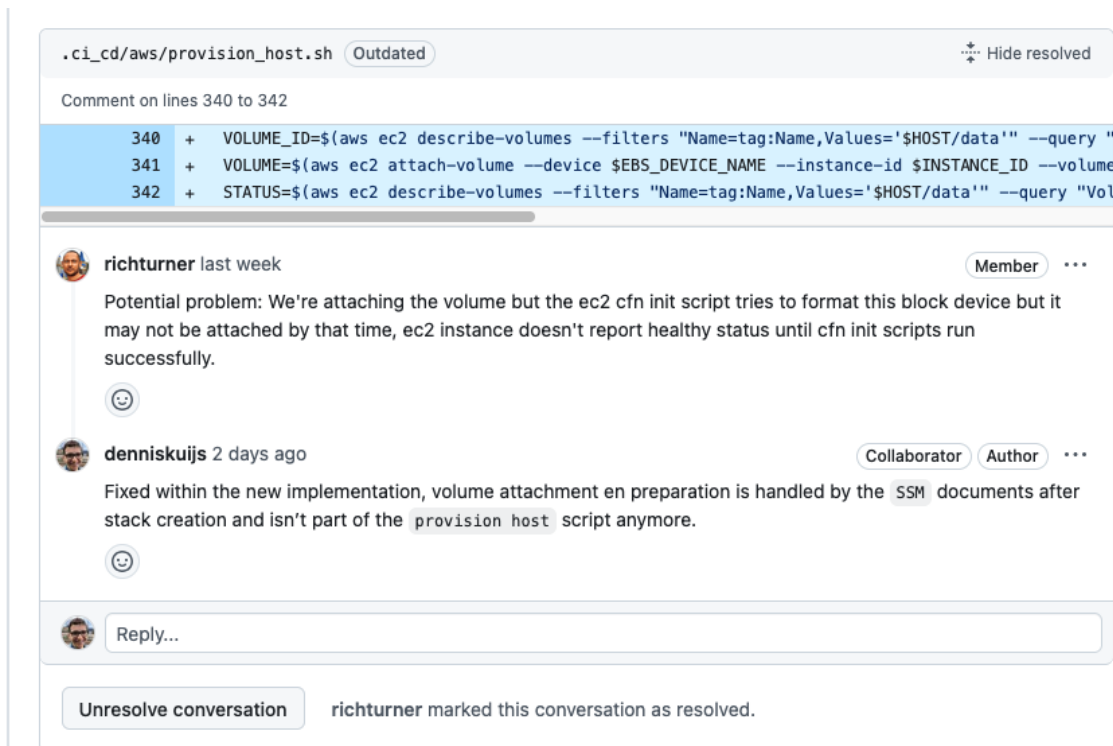


Figure 11: Rich mentioned that this solution can properly result in a potential problem

4.8. Move provisioning DLM Policy to create-ec2 stack

In my initial approach, I included logic in the `provision host` script to create the `DLM` policy using a separate `CloudFormation` stack. However, since `DLM` is a native part of `EC2`, it made more sense to move this part into the `create-ec2` `CloudFormation` stack. This change simplifies the `provision host` script and ensures the `DLM` policy is created automatically once the `EBS` data volume is successfully provisioned.

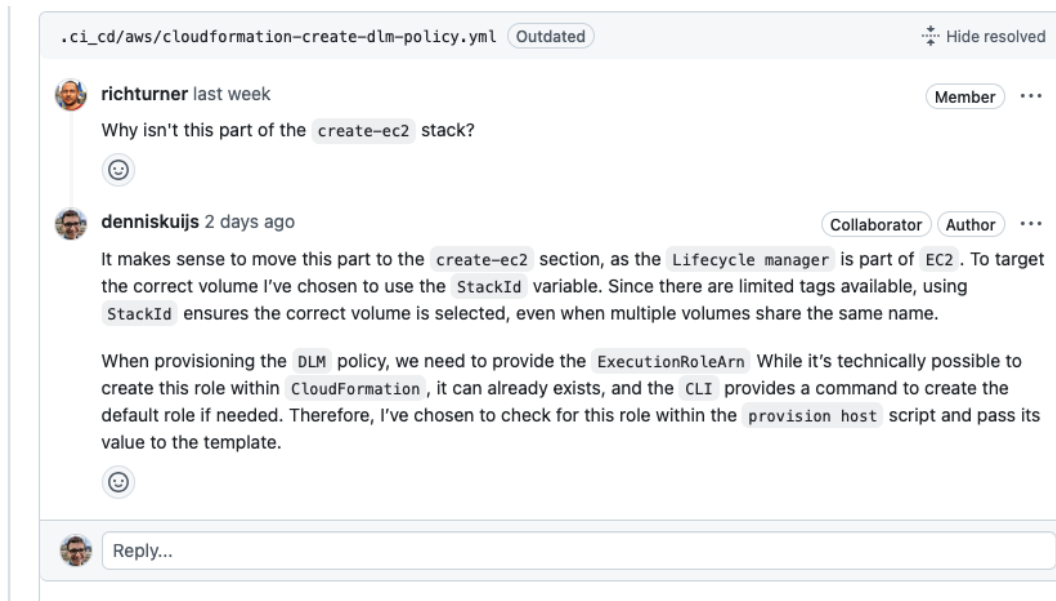


Figure 12: Rich proposed to move the DLM Policy creation to the create-ec2 CloudFormation Stack

4.9. Move provisioning EBS data volume to create-ec2 stack

Initially, I considered splitting the provisioning process to avoid the risk of the volume being detached or modified during `CloudFormation` updates.

However, after moving the volume creation into the `create-ec2` template, I found out that only the `root` volume is affected during updates, while data volumes remain unchanged. To add an extra layer of safety, I've applied `DeletionPolicy: Snapshot`, which ensures a final `snapshot` is taken before the volume is deleted. Because the volume must reside in the same `availability zone` as the instance, it is now provisioned only after the instance is fully created and the `cfn-signal` returns a `SUCCESS` status.

I've also removed the steps for creating the filesystem and mounting the volume from the `cfn-script` section. Instead, that logic is now handled via `SSM` documents, which are executed by the `provision host` script after the stack is successfully created.

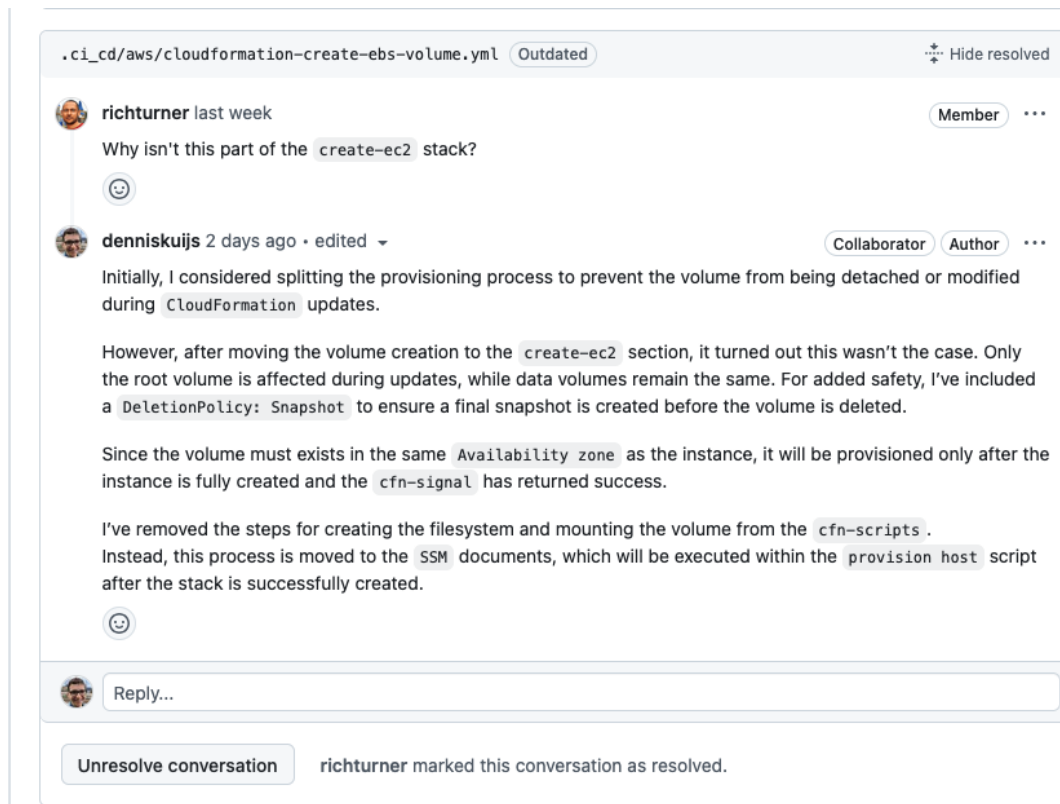


Figure 13: Rich proposed to move the EBS data volume creation to the `create-ec2` CloudFormation Stack

5. Additional changes after second feedback round

After refining my implementation based on the initial feedback, I requested another review. This time only a few minor changes were suggested.

5.1. Changing variable names and comments

To maintain consistency in naming, I replaced the slash with a hyphen in the `EBS` data volume name. This name appears in the AWS management console to make it easier to identify each volume.

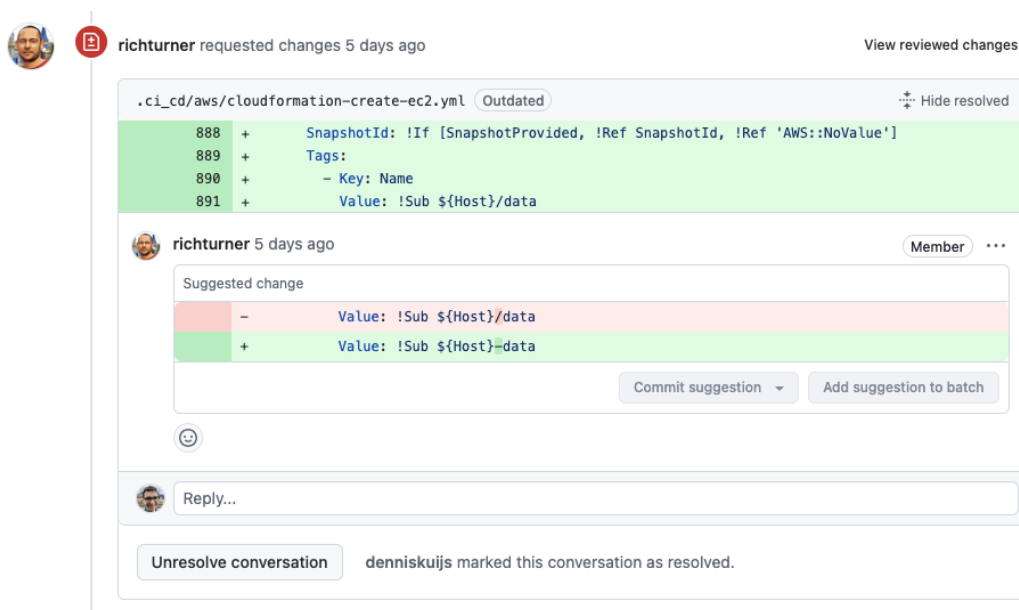


Figure 14: Updated the name of the new EBS data volume

With this name change it also needs to be updated in the `provision host` script to correctly retrieve the `VolumeId`, which is then passed to the `SSM` documents for `attaching` and `mounting` the volume.

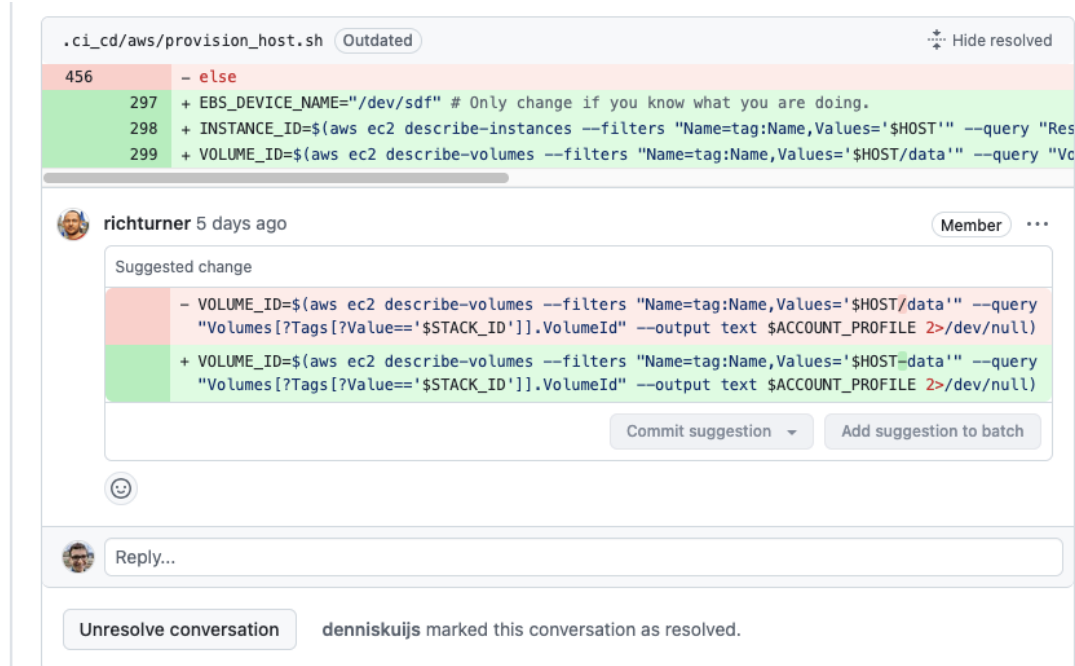


Figure 15: Updated the name of the new EBS data volume

I also corrected a few comments to fix spelling errors.

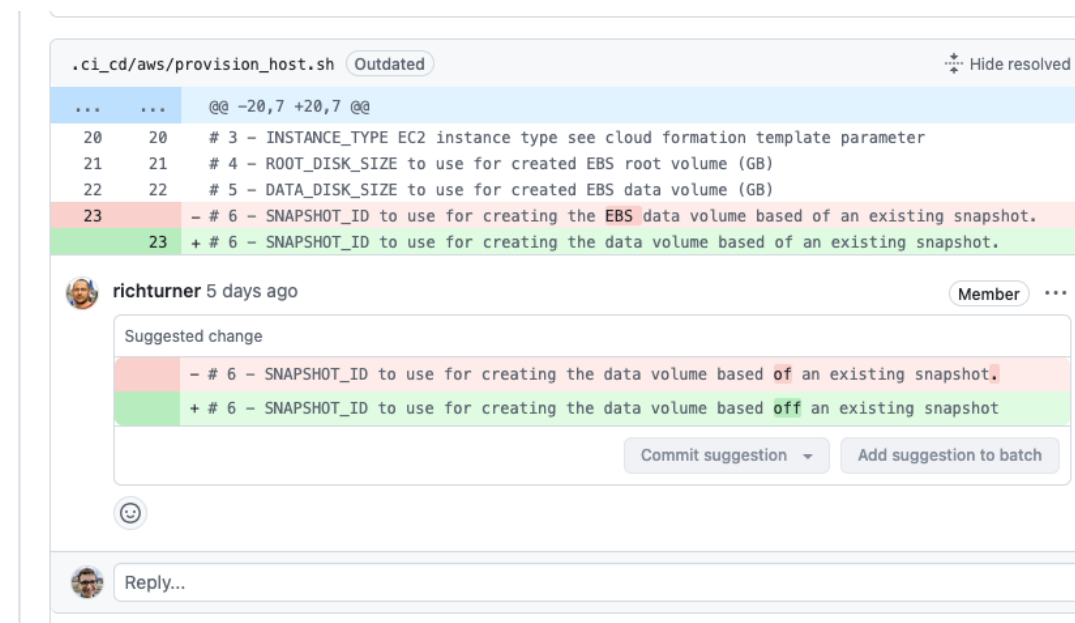


Figure 16: Correct spelling errors

5.2. Improving the filesystem check in the SSM document

I previously explained that the current implementation isn't ideal. When the volume is attached for the first time, there is no filesystem yet. which results in a `900` second wait loop before the initial filesystem is created. I've improved this by combining the volume attachment and filesystem check into a single command. Since the filesystem check depends on the volume being attached, running them separately can cause issues since these commands are running asynchronous and don't wait for each other to finish.

.ci_cd/aws/cloudformation-create-ssm-documents.yml
Outdated

Comment on lines 122 to 137

```

122 + FILESYSTEM=$(blkid -o value -s TYPE {{ DeviceName }})
123 +
124 + count=0
125 + while [[ -z "$FILESYSTEM" ]] && [ $count -lt 30 ]; do
126 +     echo "Filesystem is currently not available .. Sleeping 30 seconds"
127 +     sleep 30
128 +     FILESYSTEM=$(blkid -o value -s TYPE {{ DeviceName }})
129 +     count=$((count+1))
130 + done
131 +
132 + if [ -z "$FILESYSTEM" ]; then
133 +     mkfs -t xfs {{ DeviceName }}
134 +     mount {{ DeviceName }} /var/lib/docker/volumes
135 + else
136 +     mount {{ DeviceName }} /var/lib/docker/volumes
137 + fi

```

richturner

5 days ago

Member

...

So the first time the volume is attached we are going to wait 900 seconds before we create the filesystem. There must be a quicker way to determine if the block device has a filesystem, error codes or another command?

denniskuijs

1 minute ago

Collaborator

Author

...

Good point! I've improved the implementation by combining the volume attachment and filesystem check into a single command. Since the filesystem check depends on the volume being attached, running them separately can cause issues since these commands are asynchronous and don't wait for each other to finish.

In my last commit, the volume is first attached to the machine. Once the attachment is completed successfully, the script proceeds to retrieve the filesystem and applies the appropriate logic based on that.

- No filesystem:
Creates a new XFS filesystem
Mount volume to the /var/lib/docker/volumes directory
- Filesystem exists:
Mount volume to the /var/lib/docker/volumes directory

When a volume is attached for the first time and has no existing filesystem, the script will create one immediately after the volume is attached. Instead of waiting 900 seconds before proceeding.

Reply...

Figure 17: Rich proposed to improve the filesystem check to remove the 900 second waiting period

5.3. Replace contents of existing volume

One of my colleagues is wondering if we could replace the contents of the existing volume with a `snapshot`. Unfortunately, this is not possible as snapshots can only be specified during the initial creation of the volume. Additionally, accessing the contents of the `snapshot` directly is difficult

because snapshots are stored in a hidden `s3` bucket managed by Amazon.

The only way to retrieve data from a `snapshot` is to create a new volume from it, then attach and mount that volume to an `EC2` instance to inspect and access the data via the filesystem.

To improve flexibility, we could extend our implementation to also support provisioning a host based on an existing data volume using its `VolumeId`. This would give us three different options:

- No `SnapshotId` or `VolumeId` specified: An empty data volume is created, `attached` and `mounted`.
- `SnapshotId` specified: a new data volume is created, attached and mounted from the given `snapshot`.
- `VolumeId` specified: the existing volume is used directly, no new volume is created.

I will continue investigating the options for this in the upcoming weeks.

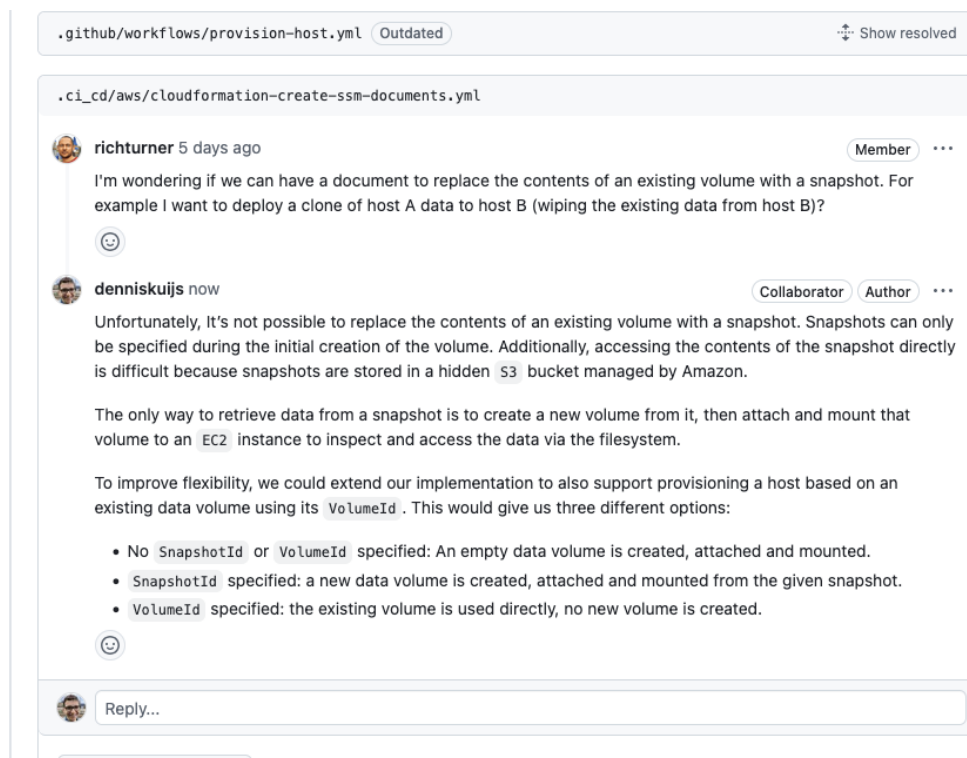


Figure 18: Investigating the option to replace an EBS data volume with an existing snapshot

6. Tests

To make sure the implementation is working as expected, I executed several tests by running the script in the `CI/CD` workflow on `GitHub Actions` .

60 workflow runs			Event ▾	Status ▾	Branch ▾	Actor ▾
✓	Provision Host Provision Host #163: Manually run by denniskuijs	master	2 days ago 7m 4s	...		
✓	Provision Host Provision Host #162: Manually run by denniskuijs	enhancement/ebs-volume-cr...	3 days ago 7m 22s	...		
✓	Provision Host Provision Host #161: Manually run by denniskuijs	enhancement/ebs-volume-cr...	3 days ago 7m 27s	...		
✓	Provision Host Provision Host #160: Manually run by denniskuijs	enhancement/ebs-volume-cr...	3 days ago 6m 48s	...		
✓	Provision Host Provision Host #159: Manually run by denniskuijs	enhancement/ebs-volume-cr...	3 days ago 7m 21s	...		
✓	Provision Host Provision Host #154: Manually run by denniskuijs	enhancement/ebs-volume-cr...	last week 21m 21s	...		
✓	Provision Host Provision Host #152: Manually run by denniskuijs	enhancement/ebs-volume-cr...	last week 7m 37s	...		
✓	Provision Host Provision Host #151: Manually run by denniskuijs	enhancement/ebs-volume-cr...	last week 6m 18s	...		
✓	Provision Host Provision Host #149: Manually run by denniskuijs	enhancement/ebs-volume-cr...	last week 7m 22s	...		
✓	Provision Host Provision Host #144: Manually run by denniskuijs	enhancement/ebs-volume-cr...	last week 7m 5s	...		
✓	Provision Host Provision Host #143: Manually run by denniskuijs	enhancement/ebs-volume-cr...	last week 6m 3s	...		

Figure 19: The provision_host CI/CD workflow has ran multiple times to test my implementation

I also performed a manual test with a fellow student by testing his changes from the OpenRemote software with my implementation.

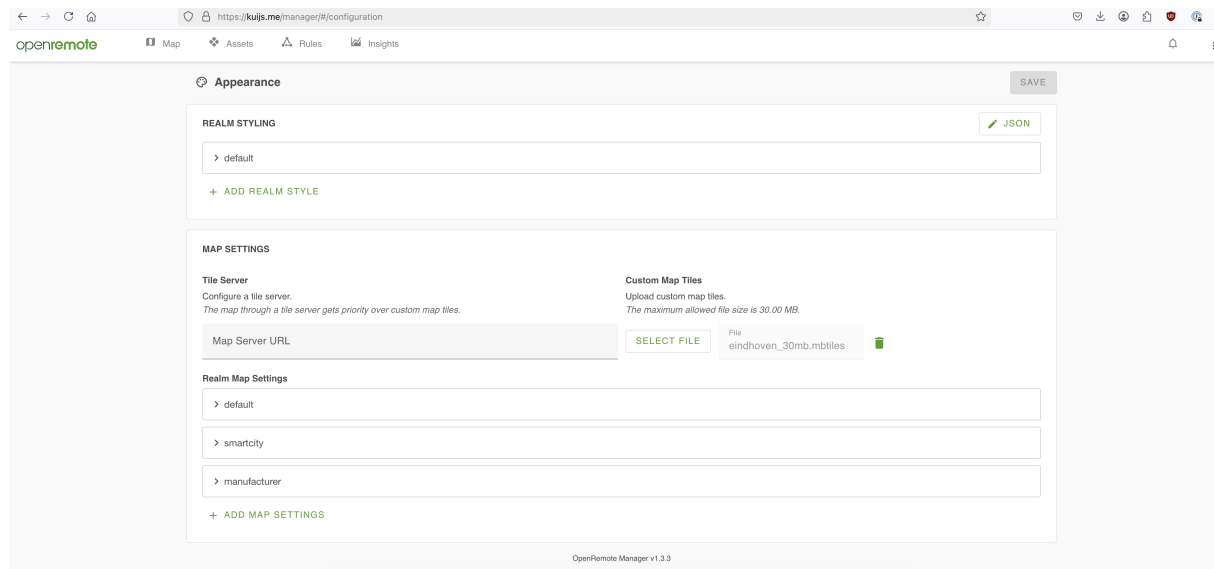


Figure 20: The map feature is visible in the OpenRemote platform

In the image above, you can see that custom `maptiles` can be uploaded through the OpenRemote Manager. These `maptiles` are stored on a `Docker` volume, which is stored on the separate `EBS` data volume.

When attaching this `EBS` data volume to another `EC2` instance, the image below shows that the `maptiles` are still available. This confirms that the implementation is working as expected. The data is correctly stored on the separate `EBS` volume and remains accessible when attached to other `EC2` instances.

```
[ec2-user@ip-172-32-1-81 ~]$ sudo lsblk -f
NAME        FSTYPE FSVER LABEL UUID                                 FSAVAIL FSUSE% MOUNTPOINTS
nvme0n1
├─nvme0n1p1 xfs     /      8891083c-63de-4431-a280-48d9dbf77c38  25.3G   15%   /
└─nvme0n1p128 vfat    FAT16  BE30-0CF9  8.6M    14%   /boot/efi
nvme1n1     ext4    1.0    8bcab1fa-c792-48ba-9c2b-99ff54463cc0  92.7G   0%    /var/lib/docker/volumes
[ec2-user@ip-172-32-1-81 ~]$ sudo ls -la /var/lib/docker/volumes/openremote_manager-data/_data/map
total 29048
drwxr-xr-x. 2 root root 4096 Apr 17 14:39 .
drwxr-xr-x. 7 root root 4096 Apr 17 14:31 ..
-rw-r--r--. 1 root root 29736960 Apr 17 14:39 eindhoven_30mb.mbtiles
[ec2-user@ip-172-32-1-81 ~]$
```

Figure 21: The `.mbtiles` file is found on the `EBS` data volume